

Report: assignment Programming Languages Ghent University, 2011

Jasper Van der Jeugt

July 5, 2011

Contents

Lazy execution: performance issues	1
Programs	1
Goldbach partition of a number	2
Quicksort	2
Hamming numbers	2
Conclusions	6
Nonogram solver	6
Choice of programming language	6
Implementation	7
Parallelization conclusions	8
Literate source code	8

Lazy execution: performance issues

Programs

We have chosen three programs to:

- calculate all Goldbach partitions of a number;
- sort a random list of numbers, using the quicksort algorithm;
- generate the hamming numbers in order.

A `Makefile` is included in the `assignment-1` directory to reproduce all benchmarks.

Goldbach partition of a number

This program (found in `goldbach-eager.oz` and `goldbach-lazy.oz`) computes the Goldbach partitions of a number `N` (which it takes as command-line argument).

We can imagine that the lazy version will be very similar to the eager version in terms of execution: in both cases, all partitions need to be generated. Hence, we may expect that the lazy version will be a little slower than the eager version, since laziness always introduces a certain overhead.

The benchmarks confirm this, we can see that the eager version is faster for all large enough inputs. For this problem, we do not seem to benefit from laziness.

The lazy function does have the benefit that we can reuse it for a program that only computes one Goldbach partition of a number. In that case, the eager function will do more work than strictly necessary.

Quicksort

`sort-eager.oz` and `sort-lazy.oz` implement a quicksort algorithm. The user passes some `N` as command-line argument to the program. The program will then generate `N` random numbers, sort them and print the result.

We have similar results compared to the Goldbach problem: all numbers need to be sorted and printed, so our algorithm does not benefit from laziness.

Note that the lazy quicksort function is, again, more composable. Imagine a scenario in which we only need the first `X` items of the sorted list (i.e. the `X` smallest items). We could use something like:

```
{List.take {QuickSort L} X R}
```

When our `QuickSort` function is lazy, the work the algorithm needs to do will be effectively reduced, since we don't need to generate anything past the first `X` elements. This is not the case for the eager version.

Hamming numbers

Our last program (see `hamming-eager.oz` and `hamming-lazy.oz`) is able to generate the first `N` hamming numbers. This `N` is passed as a command-line argument to the executable.

Here, our lazy program has two main advantages:

- it is shorter, more straightforward, and easier to understand;
- it also performs better.




Figure 1: Running time of calculating the Goldbach partitions for a given N




Figure 2: Running time of generating and sorting N numbers

Note that it is possible to write an eager program that would perform as well as the lazy program: but this would not be trivial. Basically, we would have to implement some sort of cache or memoization, and at that point we'd basically be implementing features of lazy evaluation in our eager language.

We can conclude that this simply is an algorithm which translates better to a lazy programming language than to an eager programming language.




Figure 3: Running time of generating N hamming numbers in order

Conclusions

It is obvious that the choice between lazy and eager evaluation, from the viewpoint of the programming language designer, is not an easy one. We have shown that both methods have advantages and disadvantages.

One of the key advantages I like most about lazy evaluation, is that it allows better code reuse in some cases. This was the case for the Quicksort and Goldbach problems, and another very simple example is the `{Or Xs}` function, which returns `true` if at least one of the elements in `Xs` is `true`, and `false` in all other cases. Suppose this is declared as a lazy function, and we also have a lazy `{Map Xs F}` function, which applies the function `F` to each element in `Xs`.

We can then define an `{Any Xs P}` function, by reusing the `Or` and `Map` functions:

```
declare
fun lazy {Any Xs P}
  {Or {Map Xs P}}
end
```

This reuse is possible using eager evaluation — but the result would not be optimal. In our lazy version, the (potentially expensive) predicate would not be applied to the elements in `{List.drop Xs N}` when the element at position `N` already satisfied the predicate.

On the other hand, I think a major disadvantage of lazy evaluation is that it is sometimes very hard to predict time and space complexity, in addition to the overhead lazy programs have.

If I were to design and implement a programming language, I would thus choose between:

- A lazily evaluated language with good support for eager evaluation
- An eagerly evaluated language with good support for lazy evaluation

And from those two, I would pick the former, because I personally prefer increased composability over better performance. Furthermore, most lazy programs can be made almost as fast as their eager counterparts, by performing a strictness analysis and adding a little strictness where necessary.

Nonogram solver

Choice of programming language

I chose to implement the nonogram solver in the Haskell programming language ¹. Haskell was chosen for a number of reasons:

¹<http://www.haskell.org>

- availability of a high-performance compiler, GHC ²;
- good Literate Programming support;
- semi-implicit parallelization using annotations;
- referential transparency which allows a more declarative programming model.

Implementation

The sequential and parallel nonogram solvers are implemented in the `Nonogram.lhs` file. This is a literate Haskell file ³ containing the report as well as the carefully explained source code.

The `Main.hs` file has an entry point for an executable to solve a puzzle concurrently, and a simple parser for a standard nonogram file format ⁴.

We can compile and run the program like this:

```
$ make
ghc --make -O2 -threaded Main.hs -o nonogram-solver
[1 of 1] Compiling Nonogram      ( Nonogram.lhs, Nonogram.o )
[2 of 2] Compiling Main          ( Main.hs, Main.o )
Linking nonogram-solver ...
$ ./nonogram-solver 20x20.nonogram
Parsing 20x20.nonogram
Solving 20x20.nonogram
-----XXX-----
-----XXXXX-----
-----XXX-X-----
-----XX--X-----
-----XXX-XXX-XXXX--
----XX--XX---XXXXXXX
--XXXXXXX-X---X-----
-XXXX---XX--XX-----
-----X---X-----
-----XXX--X-----
-----XXXXXX-----
-XX---XXXXXXX-----
XXXXXX--XXX-X-----
X-XX--XX-X--X-----
---XXXX--X-X--XXX---
-----XXXX-XX-XX--
-----XXX--XXX-X--
-----XXX---XXX---
-----XXX-----
-----XX-X-----
```

²<http://www.haskell.org/ghc/>

³http://www.haskell.org/haskellwiki/Literate_programming

⁴<http://www.comp.lancs.ac.uk/~ss/nonogram/fmt2>

Note that you might need to install the `parallel` package⁵, for example using `cabal-install`⁶. The `parallel` package is based on the *Algorithm + Strategy = Parallelism* paper⁷ and gives us a high-level interface to add parallelism to our program.

Parallelization conclusions

We can now compare the performance of the sequential program to the performance of the parallel program. We use the excellent criterion library⁸, aimed at benchmarking Haskell code. The code we used to benchmark the the programs is located in the `Benchmark.hs` file, along with some sample puzzles. To reproduce the benchmarks, the `Makefile` contains the `benchmark-sequential` and `benchmark-parallel` targets, which benchmark the sequential and parallel program.

We ran these benchmarks on an Intel(R) Core(TM)2 Duo CPU E8400 @ 3.00GHz. Because we use a dual-core computer, we can at most expect a 100% speedup (i.e. a reduction of the running time by 50%).

	5x5	10x10	15x15	20x20
sequential	10.10453 us	202.7505 us	5.911520 ms	250.8528 ms
parallel	22.05260 us	392.9396 us	4.348415 ms	155.6637 ms

Looking at the results, parallelization introduces a large overhead for the smaller puzzles, but gives us an advantage for larger puzzles. Because of our implementation of the branch-and-bound-based algorithm, we know that for each branch, a `par` call is made. However, this is only useful if the branch is a non-trivial computation: otherwise, the overhead of the `par` call outweighs the benefits of parallelization. We can conclude that the parallelization for this algorithm is only useful for large enough input sets: otherwise, the overhead outweighs the speedup.

For the 20x20 puzzle, we see that the running time is reduced by 37.9461%. This is not the 50% reduction we were naively hoping for, but it is not a bad result, since the parallelization of the program was almost trivial (just replacing the branching function). We also have to keep in mind that not everything can happen in parallel: e.g. joining the results of two branches happens on one core.

Literate source code

Here, we give the full source code to the programs, annotated in Literate Programming style.

```
module Nonogram
```

⁵<http://hackage.haskell.org/package/parallel>

⁶<http://hackage.haskell.org/trac/hackage/wiki/CabalInstall>

⁷Philip W. Trinder, Kevin Hammond, Hans-Wolfgang Loidl, and Simon L. Peyton Jones. Journal of Functional Programming, 8(1), Jan 1998. <http://www.macs.hw.ac.uk/~dsg/gph/papers/abstracts/strategies.html>

⁸<http://www.serpentine.com/blog/2009/09/29/criterion-a-new-benchmarking-library-for-haskell/>




Figure 4: Performance comparison of the nonogram solver (log scale on y axis)

```

    ( Description
    , sequentialNonogram
    , parallelNonogram
    , putNonogram
    ) where

import Control.Monad (when, mplus, foldM)
import Control.Parallel (par, pseq)
import Data.IntMap (IntMap)
import qualified Data.IntMap as IM

```

A nonogram grid is built from cells in two possible colors. For representing the value of a cell, we use a simple datatype called `Color`.

```

data Color = White | Black
           deriving (Show, Eq)

```

A nonogram is then defined as a matrix of colors:

```

type Nonogram = [[Color]]

```

The input is a list of natural numbers for every row and column. We call such a list a `Description` and represent it as an `Int` list:

```

type Description = [Int]

```

We use an algorithm that attempts to solve the puzzle top-down, i.e. starting at the top row and finishing at the bottom row. It is implemented as a branch-and-bound algorithm, searching the solution space.

As we solve more and more rows, we also learn new, partial information about the columns — in addition to the `Description` of the columns. We use this to “bound” in certain cases, if the current search space does not match our partial information.

```

data Partial = MustBe Color
             | BlackArea Int
             deriving (Show, Eq)

```

Initially, this partial information is deduced from the descriptions alone. Hence, we have a function to convert a `Description` into `Partial` information.

```

fromDescription :: Description -> [Partial]
fromDescription = map BlackArea

```

When we reach the bottom of the nonogram, we will have a “final state” of the partials for every column. We then need to check if we “used up” all black areas: this function checks if the partial information can be considered empty. Note that we do allow **White** fields in empty partials (since they do not add to the description counts).

```
emptyPartial :: [Partial] -> Bool
emptyPartial [] = True
emptyPartial (MustBe Black : _) = False
emptyPartial (BlackArea _ : _) = False
emptyPartial (_ : xs) = emptyPartial xs
```

We need to keep a `[Partial]` for every column. We can’t store them in a list because we want fast random access. We use an `IntMap`, a purely functional tree-like structure (big-endian patricia trees) which performs fast for lookups and insertions.

```
type Partials = IntMap [Partial]
```

The following function adds a cell to partial information. It returns a value in the `Maybe` monad:

- if the cell is inconsistent with previous information, it returns `Nothing`;
- otherwise, it returns the updated partial information.

On update, it might consume or add to the partial information. Note that the partial information always represents “what comes next” in the column, so we only need to inspect/modify the first few elements of the list.

```
learnCell :: Color -> [Partial] -> Maybe [Partial]
learnCell White [] = Just []
learnCell Black [] = Nothing
learnCell y (MustBe x : ds) = if x == y then Just ds else Nothing
learnCell White ds = Just ds
learnCell Black (BlackArea n : ds) = Just $
    replicate (n - 1) (MustBe Black) ++ (MustBe White : ds)
```

We also provide a convenience function to call `learnCell` for a particular column (specified by its 0-based index).

```
learnCellAt :: Color -> Partials -> Int -> Maybe Partials
learnCellAt cell partials index = do
    ds <- learnCell cell $ partials IM.! index
    return $ IM.insert index ds partials
```

In Haskell, we can abstract over pretty much anything. We need to write two versions of our nonogram solving algorithm — a simple one and a concurrent one. Since we want as little code duplication as possible, we can do this by abstracting over the way we deal with branches in our search tree.

More concrete, we have a branching strategy (`Branching`) which decides how two branches which might or might not yield a result (the `Maybe a`'s) are composed. We also provide sequential and parallel branching:

```
type Branching a = Maybe a -> Maybe a -> Maybe a
```

```
sequential :: Branching a
sequential = mplus
```

```
parallel :: Branching a
parallel x y = x 'par' y 'pseq' mplus x y
```

Parallelization of Haskell programs can be easily done using `par` annotations, which attempts to spark the calculation of a thunk on a free core (if possible). While `par` is a very cheap function call, it is not *free*. Therefore, one always has to make sure the computations we spark on other cores are big enough, so we minimize the overhead caused by `par`.

This implies it could be a better strategy to only use parallel branching if we are in the upper part of the search tree (i.e. depth is less than a given `n`). However, after experimenting with this, this `n` was highly dependent of the used input set, and the speedup was only marginally better than the simpler solution using just `parallel` — so I decided not to use this more advanced strategy.

The strategy is used in the `solve` function, which holds most of the main solver logic.

The `solve` function is actually a wrapper around a `solve'` function which does the actual work. This is an optimization called the static argument transformation ⁹.

```
solve :: Branching Nonogram -> Int -> Int -> [Description] -> Partial
      -> Maybe Nonogram
solve branch width = solve'
  where
    solve' column descriptions partials
```

There are a number of cases that need to be considered. First, suppose there are no more descriptions for the rows. In this case, we know that the knowledge we have about the column partials must be empty: we can no longer place any black cells. If this is the case, we have a correct solution (the empty one), otherwise, our solver fails by returning `Nothing`.

```
    | null descriptions =
      if all emptyPartial (map snd $ IM.toList partials)
      then return [[]]
      else Nothing
```

⁹<http://blog.johantibell.com/2010/09/static-argument-transformation.html>

If we have at least one row description, we check to see if that row description is empty. Suppose this is the case. This means no more black areas should be placed in this row, so we just fill it up with white cells. We then recursively call `solve'` to solve the other rows.

```
| null rd = do
  ps <- foldM (learnCellAt White) partials [column .. width - 1]
  rows <- solve' 0 rds ps
  return $ replicate (width - column) White : rows
```

Since we know now that the row description is not empty, we have at least one black area on this row. We consider two possibilities:

- the black areas starts at the beginning of the row (consider the beginning of the row as indicated by the `column` argument);
- we have at least one white cell, followed by a row with the same black areas.

This is where our search tree branches between the two cases, `branch` and `skip`, defined later.

```
| otherwise = branch place skip
```

Some definitions of previously used values:

```
where
  (rd : rds) = descriptions
  (l : ds) = rd
```

The case where the black area is at the beginning at the row has a pretty long definition but the logic behind the code is actually quite straightforward:

- we fail if the black area cannot be placed due to the fact there are insufficient cells left;
- we update the the column partial knowledge with these new black cells;
- we place a white cell after the black area, two because black areas cannot be contiguous (if the end of the black area touches the border of the grid, we can skip this);
- we recursively call `solve'` to solve the rest of the row, then add the cells we just calculated to the returned solution.

```
place = do
  when (column + 1 > width) Nothing
  ps <- foldM (learnCellAt Black) partials
    [column .. column + 1 - 1]

  let atEnd = column + 1 == width
```

```

ps' <- if atEnd then return ps
      else learnCellAt White ps (column + 1)

(row : rows) <- solve' (column + 1 + 1) (ds : rds) ps'
let row' = if atEnd then row else White : row
return $ (replicate 1 Black ++ row') : rows

```

If we choose not to place the black area at the beginning of the row, we just need to add a white cell and recursively call `solve'` again. This case can fail as well — when we've reached the far right side of the grid, we can no longer place white cells.

```

skip = do
  when (column >= width) Nothing
  ps <- learnCellAt White partials column
  (row : rows) <- solve' (column + 1) ((1 : ds) : rds) ps
  return $ ((White : row)) : rows

```

The `nonogram` function helps us in converting the column descriptions into the partial information we need, and thus provides a nicer interface to the programmer than `solve` does.

```

nonogram :: Branching Nonogram -> [Description] -> [Description]
          -> Maybe Nonogram
nonogram branch rows columns = solve branch (length columns) 0 rows state
  where
    state = IM.fromList (zip [0 ..] $ map fromDescription columns)

```

We provide a sequential and a parallel program:

```

sequentialNonogram :: [Description] -> [Description] -> Maybe Nonogram
sequentialNonogram = nonogram sequential

```

```

parallelNonogram :: [Description] -> [Description] -> Maybe Nonogram
parallelNonogram = nonogram parallel

```

At last, the `putNonogram` function allows us to print a solution we found (of the type `Maybe Nonogram`) to standard output.

```

putNonogram :: Maybe Nonogram -> IO ()
putNonogram Nothing = putStrLn "No solution found"
putNonogram (Just s) = mapM_ (putStrLn . concatMap showCell) s
  where
    showCell Black = "X"
    showCell White = "-"

```